Rivero

WHITE PAPER

Disputes in the age of agentic commerce

Why issuer dispute operations will feel the shift before most teams have a framework for it

INSIDE THIS WHITE PAPER

Contents

ABOUT THIS WHITE PAPER

Payments are becoming autonomous, continuous, and increasingly machine-speed. Dispute operations are not. This paper maps the structural break ahead, and what an agentic-ready back office looks like.

EXECUTIVE SUMMARY

The dispute model wasn't built for machine speed

Card payments are becoming autonomous, with AI agents enabling machine-speed execution. The back-office operations that handle them are not. As AI agents begin to transact on people's behalf, the gap between the two is about to widen sharply, and disputes are where it shows first.

At the same time, first-party fraud has industrialised and merchants are countering it with AI. Issuers are caught in the middle, and the manual, staff-driven dispute model simply can't absorb what is coming.

324M projected global chargebacks by 2028	**$10.32** to manually process a single dispute	**84%** of banking leaders see AI agents as the greatest fraud threat	**70%** of the dispute process Amiko can automate

THE STRUCTURAL BREAK

Payments are moving at machine speed while dispute handling remains at human speed. Issuers can't hire their way out of that asymmetry. The gap comes down to a strategic decision about how a bank scales.

This white paper maps the shift ahead: the new categories of dispute, the economics that break first, and the operating model that keeps pace.

THE AGENTIC RESPONSE

DEFLECT • AUTOMATE • RESOLVE

The structural break in payment operations

The convenience of the digital world lets us manage our lives online, but it has also made it easier for cardholders to dispute a transaction with a single click. Global chargeback volumes are projected to reach 324 million by 2028, a double-digit increase over the 261 million forecast for 2025. Correspondingly, global card fraud losses are projected to reach $41.06 billion by 2030.

Against this backdrop of rising volumes, a new structural asymmetry is emerging. The industry is rapidly deploying infrastructure for AI-initiated transactions, enabling "Commerce 3.0" where AI agents act and transact on behalf of consumers. Payments are becoming autonomous, continuous, and machine-speed.

However, the back-office operations that handle exceptions remain reliant on human-speed processes, caseworkers manually interpreting claims, gathering evidence and navigating legacy systems. This linear, manual model costs up to $10.32 per dispute. When high-frequency, machine-initiated micro-transactions meet manual exception handling, the unit economics of issuer operations will collapse.

THE VOLUME
Global chargeback volume

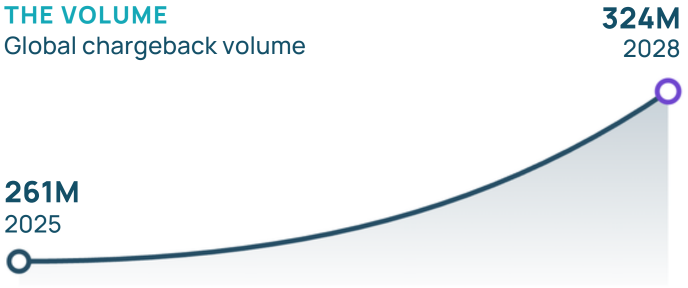

THE COST

$41.06B

Projected global card fraud losses

$10.32 to manually process a single disputed transaction, the cost that collapses at machine speed.

Furthermore, fraudsters and merchants have already industrialised their operations. Retail refund fraud now operates as a "crime-as-a-service" economy, and merchants are using AI to fight chargebacks with compelling evidence. Issuers are caught in the middle.

THE BOTTOM LINE

To survive the coming wave of disputes, issuers can't simply expand manual operations. They must adopt autonomous, agentic dispute platforms built for intelligent intake, zero-touch resolution, and programmatic evidence processing.

SECTION 01

The agentic commerce reality

In June 2026, the shift toward autonomous purchasing reached a definitive inflection point with two major network announcements.

Visa × OpenAI

Visa is integrating Visa Intelligent Commerce directly into ChatGPT, letting AI agents execute payments with tokenised Visa credentials inside defined permissions, spending limits, merchant categories.

Mastercard AP4M

Mastercard launches Agent Pay for Machines to unlock super-fast, always-on payments. It orchestrates, and settles machine-to-machine transactions at very low latency, down to fractions of a cent.

> **AI is positioned to transform commerce more profoundly than the internet or mobile technology ever did.**
>
> – Visa, source

However, as agents take over transaction routing and real-time decision-making, new challenges arise for issuing banks. The networks anticipate friction: Mastercard's Agent Pay protocol even builds in interface standards to handle unfamiliar transactions. But while the front-end rails are fortified with tokenisation and multi-factor authentication, back-end exception handling is largely unprepared.

84%

of banking leaders now identify AI agents as the industry's greatest exploitable vulnerability for fraud. When an AI agent authorises a purchase later disputed by the customer, that dispute flows into queues entirely unequipped for machine-to-machine conflict resolution.

SECTION 02

The industrialisation of first-party fraud

Before agentic commerce fully scales, issuer back offices are already collapsing under the weight of "friendly fraud" and first-party misuse. First-party fraud and third-party fraudulent chargebacks account for roughly 45% of merchant chargeback volume globally.

The normalisation of this fraud is a growing crisis. What was once considered opportunistic consumer behaviour has professionalised into a highly structured "crime-as-a-service" underground economy. Recent research from the University of Portsmouth, Mapping the online economy of refund fraud, reveals a sophisticated network of an estimated 60,000 to 90,000 participants.

45%	60-90K	$2,419
of merchant chargeback volume is first- or third-party fraud	estimated participants in the refund-fraud economy	mean fraudulent claim value in sampled networks

Operating as an organised crime-as-a-service model, professional refund vendors advertise directly to consumers, promising to secure illicit refunds in exchange for a fee that typically ranges between 13% and 30% of the refunded amount. These illicit networks meticulously share intelligence on the most vulnerable digital storefronts, continuously evolving their social engineering tactics and manipulation methods to systematically exploit loopholes in specific retailer return policies. This directly impacts issuing banks.

WHY THIS LANDS ON ISSUERS

When a merchant detects and rejects a fraudulent refund, the consumer, often coached by the vendor, simply pivots and files a chargeback. Because traditional intake forms can't distinguish a genuine delivery failure from a coached fraudster reading a script, the issuer processes the claim, absorbs the cost, and skews its own fraud data.

SECTION 03

A new taxonomy of disputes

Agentic commerce will exacerbate this already strained ecosystem by introducing entirely new dispute categories that break existing reason codes. When an AI acts on a consumer's behalf, the classic fraud-versus-friendly-fraud distinction blurs.

INTENT: ‹ Friendly — Malicious ›

INITIATOR: Human — Agent

INITIATOR / INTENT	‹ Friendly	Malicious ›
Human	**Unrecognised agent transactions** Cardholders fail to recognise machine-initiated micro-charges on their statement and file friendly-fraud claims via banking apps.	**Industrialised refund fraud** Crime-as-a-service retail fraud, increasingly fought by merchants armed with AI-driven compelling evidence.
Agent	**Machine-to-machine exceptions** High-frequency, sub-cent charges that fail settlement or duplicate, creating massive volumes of technical exceptions.	**Compromised-agent & mandate-exceeded** Hijacked AI assistants transacting with tokenised credentials, or agents acting beyond their limits.

WHY THIS LANDS ON ISSUERS

Issuers already struggle to categorise existing claims, often labelling every chargeback as fraud when no suitable code exists. Processing agent-initiated disputes using a traditional taxonomy will lead to misclassification, inaccurate fraud reporting, and increased scheme penalties.

SECTION 04

The economics break first

The traditional model runs on a linear relationship between transaction volume and back-office operational capacity. This model is economically unsustainable in an agentic world.

$9-10.32 processing overhead per disputed case Source	**10-45 min** manual case preparation per dispute Source	**200+** back-office staff at some US banks solely for chargebacks Source

Financial institutions need roughly one full-time employee for every $13,000-$14,000 in incoming annual cardholder disputes. Faced with the high-velocity, low-value exceptions of machine-speed commerce, that cost structure simply collapses.

THE WRITE-OFF TRAP

$0.50 disputed micro-transaction vs. **$10** cost to investigate it

When investigation costs more than the claim, issuers are forced to write off, and the average write-off is already $17.90.

While writing off low-value claims saves immediate processing time, it inadvertently trains opportunistic fraudsters and serial disputers to repeatedly exploit the system, compounding long-term losses.

SECTION 05

The data foundation – from raw strings to resolution

Before issuers can automate disputes, they must fix the underlying data. Whether a customer is checking their feed or filing a claim, everything starts with the transaction data left behind, and too often that data is fragmented, messy, or incomplete.

RAW STRING • DIGITAL NOISE

POS 8273 *SKYLINE APPLE.COM/BILL

Confusing strings fuel unnecessary disputes, which is a leading cause of "friendly fraud" driven purely by confusion.

→

ENRICHED DATA • AN ASSET

Sky Airlines Travel – Zürich, CH

Clean merchant name, official logo, accurate category and exact GPS location of the purchase.

To build an automated, agentic dispute response, issuers require enriched transaction data. Modern enrichment platforms clean transaction feeds by accurately identifying the merchant name, providing the official logo, assigning proper categorisation, and pinpointing the exact GPS location of the purchase.

When a cardholder sees a clear, instantly recognisable transaction, complete with the merchant's logo and exact location, the likelihood of them filing a false "unrecognised transaction" claim drops dramatically.

KEY TAKEAWAY

Clean data is a UX enhancement and the prerequisite infrastructure for dispute deflection and intelligent automation.

SECTION 06

Evidence has changed sides, and liability shifts

In traditional card disputes, the investigation relies heavily on interpreting a cardholder's subjective narrative. In agentic commerce, the nature of evidence changes entirely.

Agentic transactions generate highly structured, machine-readable data trails. This includes agent identity IDs, tokenised credentials, explicit consent mandates, and specific permission scopes outlining the exact instructions the consumer gave to their AI assistant.

63%

of merchants plan to increase spending on fraud tools and technology (2025 Global eCommerce Payments & Fraud Report, Visa Acceptance Solutions, Verifi & Merchant Risk Council), shifting toward digital screening and automated rule updates for compelling evidence. Merchants now fight back against first-party misuse at scale.

Consequently, the dispute investigation shifts from interpreting a human story to evaluating a programmatic mandate trail. Human caseworkers are poorly equipped to rapidly ingest, parse, and verify complex API logs, cryptographic consent tokens, and merchant compelling evidence. Conversely, AI-driven platforms excel at processing this structured data. The dispute operations function that can programmatically consume and analyse agentic evidence will achieve accurate, rapid resolutions.

87%

of merchants now submit compelling evidence to respond to first-party misuse disputes, up from 83% the previous year. Merchants are submitting more data points than ever, including IP addresses, user account/login IDs, and delivery information, turning structured transaction trails into a frontline defence against first-party misuse at scale.

SECTION 07

Agentic disputes need agentic dispute management

The defining asymmetry of the coming decade is machine-speed payments met by human-speed exception handling. The only back-office operation capable of keeping pace with autonomous commerce is an autonomous dispute management platform. Industry front-runners are abandoning manual intake and adopting agentic dispute management. By deploying intelligent automation, banks can separate genuine fraud from first-party misuse, process structured machine evidence, and resolve claims with zero-touch efficiency.

01 Intelligent virtual agents at intake

To tackle volume volatility and combat coached refund fraudsters, issuers use 24/7 Virtual Agents integrated directly into mobile banking apps. When a cardholder initiates a claim, the Virtual Agent uses enriched transaction data to ask the right questions, enrich context, and challenge suspicious claims. By helping cardholders accurately recognise transactions through clean data, and by identifying the subtle patterns of first-party misuse, this intelligent layer prevents unnecessary cases from entering the back office entirely.

02 Zero-touch recovery

Instead of immediately filing a formal chargeback, agentic platforms instantly query merchant collaboration networks like Ethoca and Verifi. If a claim is valid, these integrations secure a direct merchant credit to the cardholder within the strict regulatory timelines of PSD2, often within 24 hours. Around 41% of disputes can be resolved via this zero-touch recovery method, returning funds without a human analyst ever opening the case.

03 Autopilot, Copilot, and bulk processing

For cases that proceed to the back office, rule-based automation (Autopilot) executes routine decisions based on issuer-defined risk thresholds. For complex exceptions requiring human judgment, AI-assisted Copilot surfaces intelligence directly in the agent's workflow, such as detecting serial disputer anomalies. For major fraud events, bulk actions process 50 transactions in 5-10 minutes, down from three hours.

REAL-WORLD RESULTS

Rivero's agentic dispute management tool, in numbers

Recent analysis of the Amiko platform shows the immediate impact of autonomous dispute management on real issuer portfolios.

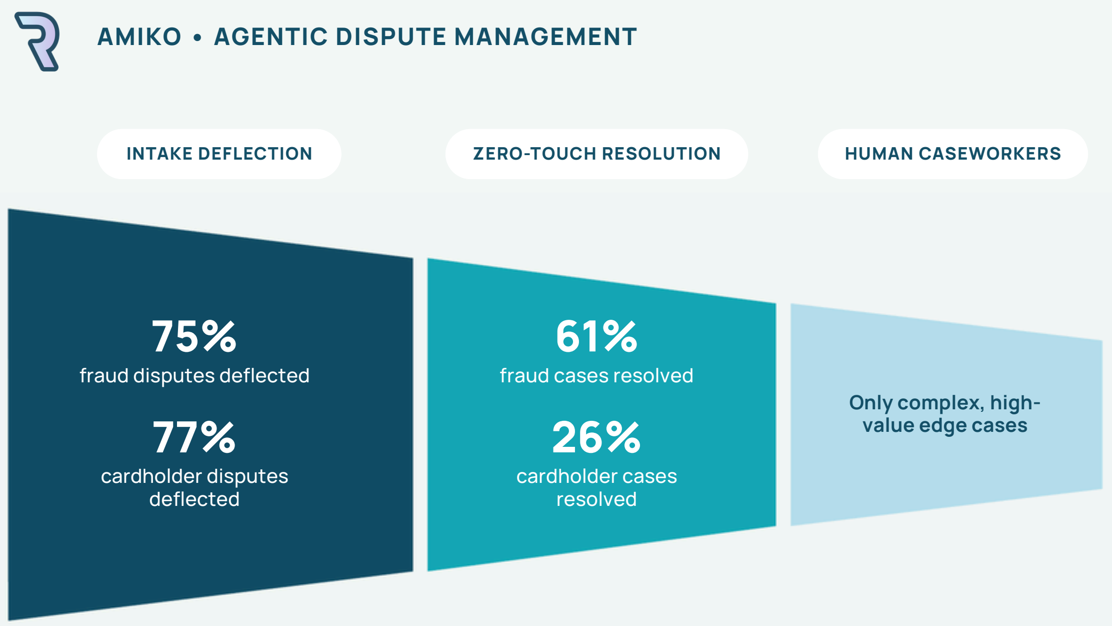

Each stage strips volume from the queue, so only the cases that truly need human judgement ever reach an analyst.

As Amiko's Autopilot and Copilot capabilities expand, these zero-touch resolution rates break the linear link between transaction volume and back-office operational capacity, which is the defining constraint that makes traditional operations unsustainable at machine speed.

Becoming agentic-ready

Issuers can navigate the evolving landscape of agentic commerce with a dispute operation that adapts at the same speed. Modernising the tech stack and deploying an intelligent, automated platform breaks the linear relationship between portfolio growth and back-office costs, transforming a regulatory burden into a highly efficient customer experience advantage.

SELF-ASSESSMENT

The agentic readiness checklist

As you evaluate your current payment operations, ask your team the following six questions.

01

Is your transaction data clear?

Do you rely on raw, confusing billing strings, or do you use enriched data (logos, clear names) to actively prevent cardholder confusion before a claim is filed?

02

Can you ingest machine-readable evidence?

Can your system parse agent identity, tokenised credentials, and merchant compelling evidence programmatically?

03

What is your true cost per dispute?

Are you trapped in a cycle of writing off low-value claims because your manual handling cost exceeds $10 per case?

04

Are you deflecting claims at intake?

Can you proactively educate cardholders and deflect the sophisticated first-party misuse tactics taught by online fraudsters?

05

Are you leveraging pre-dispute networks?

Do you automatically submit claims to merchant collaboration networks (Ethoca / Verifi) to secure zero-touch refunds within 24 hours?

06

How quickly can you adapt to scheme changes?

When Visa or Mastercard introduce new liability rules for agentic transactions, does it require a major IT project, or can you update your workflow in days?

Rivero

AGENTIC DISPUTE MANAGEMENT

See intelligent payment infrastructure in action

Discover how Rivero's Amiko platform can automate over 70% of your dispute process, deflect friendly fraud, and guarantee regulatory compliance.

Watch a demo → rivero.tech/amiko